

Overview

(As of April 28, 2024) I'm a senior researcher in the PROSE team at Microsoft, where I work on using machine learning to develop state-of-the-art program synthesis systems to make writing, fixing, and using software easier and more enjoyable. I have also worked at Facebook, Cloudera, and (in prior life) Morgan Stanley.

Education

Academic Qualifications

- 2016-2021 **PhD in Computer Science**
Massachusetts Institute of Technology, Cambridge, MA.
- 2013-2016 **Masters in Computer Science**
New York University: Courant Institute of Mathematical Sciences, NY, NY.
GPA: 3.89, MS Research/Thesis Fellowship Award Fall 2015, funding work on A2Q (an order-aware optimizing query compiler for AQuery)
- 2007-2011 **Bachelor of Arts in Economics and Minor in German Studies**
University of Pennsylvania, Philadelphia, PA.
GPA: 3.93, Phi Beta Kappa, Summa Cum Laude, Dean's List (08, 09, 10)

Industry Work Experience

- 06/2022 – **Senior Researcher PROSE Team**, Microsoft, Remote.
current
- 06/2021 – **Researcher PROSE Team**, Microsoft, Remote.
06/2022
 - Working on program synthesis technologies for a variety of developer, data scientist, and end-user applications. A lot of my work focuses on developing and applying large language models to programming tasks, such as program repair and natural language to code synthesis. As part of my job, I also manage and mentor junior researchers through the PROSE research fellowship program.
- Summer 2020 **Intern Facebook AI Research**, Facebook, Remote.
 - Worked with the SysML team at FAIR on a novel tensor compiler, writing C++ for JIT compilation, benchmarking against Halide/TVM
- Fall 2018 **Part-Time Research Visitor Big Code Team**, Facebook, Remote.
 - Applied deep learning to identify and highlight core code functionality in early ML4Code models.
- Summer 2018 **Intern Software Engineering**, Facebook, Boston.
 - Applied deep learning to code search and contributed to some of the earliest ML4Code models in this space.
- Summer 2015 **Intern Data Science**, Cloudera, San Francisco.
- 2011 – 2014 **Full-Time Securitized Credit Research Associate Non-Agency Mortgages and US Housing**, Morgan Stanley, New York.

- Summer 2010 **Richard B. Fisher Scholar** *Fixed Income Generalist Sales and Fixed Income Credit Strategy*, Morgan Stanley, New York.
- Summer 2009 **Douglas Paul Scholar** *Investment Banking and Alternative Investments*, Morgan Stanley, New York.

Academic Work Experience

- Fall and Spring 2021 **Advanced Undergraduate Research Class TA**, MIT.
- 2015 – 2016 **Graduate Course in Compiler Construction Grader**, NYU.
- Fall 2014 **Graduate Course in Programming Languages Teaching Assistant**, NYU.

Language skills

- **Programming Languages:** Proficient in: Python, Javascript/Typescript, R, C#.
- **Natural Languages:** Native fluency in English and Spanish. Working proficiency in German.

Service

- **Program Committee ICSE 2024**
- **Program Committee Table Representation Learning Workshop (at NeurIPS) 2023**
- **Program Committee Table Representation Learning Workshop (at NeurIPS) 2022**
- **Artifact Evaluation Committee OOPSLA 2020**
- **Artifact Evaluation Committee CAV 2020**
- **Artifact Evaluation Committee PPOPP 2018**

Mentoring/Advising

- Jennifer McCleary (MIT) MEng Thesis: pancreatic cancer risk modeling (Fall 2019 - January 2020)
- Alex Berg (MIT) Undergraduate research: pancreatic cancer risk modeling (Summer 2020)
- Thomas Xiong (MIT) MEng Thesis: pancreatic cancer risk modeling (Fall 2020 - Spring 2021)
- Lori Zhang (MIT) Undergraduate research: pancreatic cancer risk modeling (Summer 2020 - Spring 2021)
- Harshit Joshi (Microsoft): PROSE Research fellow, automated program repair (Fall 2021 to July 2023 – joining Stanford PhD program 2023)
- Mukul Singh (Microsoft): PROSE Research fellow, NL-to-Code (Spring 2022 to date)
- Abishai Ebenezer (Microsoft): PROSE Research fellow, automated program repair (Fall 2022 to July 2023)
- Jialu Zhang (Yale/Microsoft): Summer intern in the PROSE team, working on automated program repair (Summer 2022). Part of thesis committee.

Publications

- [1] S. Barke, C. Poelitz, C. S. Negreanu, B. Zorn, J. Cambronero, A. D. Gordon, V. Le, E. Nouri, N. Polikarpova, A. Sarkar, et al. Solving data-centric tasks using large language models. *arXiv preprint arXiv:2402.11734 (to appear NAACL 2024)*, 2024.
- [2] A. Singha, J. Cambronero, S. Gulwani, V. Le, and C. Parnin. Tabular representation, noisy operators, and impacts on table structure understanding tasks in llms. *arXiv preprint arXiv:2310.10358 (Table Representation Learning at NeurIPS 2023)*, 2023.
- [3] M. Singh, J. C. Sánchez, S. Gulwani, V. Le, C. Negreanu, M. Raza, and G. Verbruggen. Cornet: Learning table formatting rules by example. *Proc. VLDB Endow.*, 16(10):2632–2644, jun 2023.
- [4] M. Singh, J. Cambronero, S. Gulwani, V. Le, and G. Verbruggen. Emfore: Online learning of email folder classification rules. In *Proceedings of the 32nd ACM International Conference on Information*

- and Knowledge Management, CIKM '23, page 2280–2290, New York, NY, USA, 2023. Association for Computing Machinery.
- [5] M. Singh, J. Cambroner, S. Gulwani, V. Le, C. Negreanu, and G. Verbruggen. CodeFusion: A pre-trained diffusion model for code generation. In H. Bouamor, J. Pino, and K. Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 11697–11708, Singapore, Dec. 2023. Association for Computational Linguistics.
 - [6] M. Singh, J. Cambroner, S. Gulwani, V. Le, C. Negreanu, E. Nouri, M. Raza, and G. Verbruggen. Format5: Abstention and examples for conditional table formatting with natural language. *arXiv preprint arXiv:2310.17306 (to appear VLDB 2024)*, 2023.
 - [7] T. Phung, V.-A. Pădurean, J. Cambroner, S. Gulwani, T. Kohn, R. Majumdar, A. Singla, and G. Soares. Generative ai for programming education: Benchmarking chatgpt, gpt-4, and human tutors. In *Proceedings of the 2023 ACM Conference on International Computing Education Research - Volume 2*, ICER '23, page 41–42, New York, NY, USA, 2023. Association for Computing Machinery.
 - [8] T. Phung, J. Cambroner, S. Gulwani, T. Kohn, R. Majumdar, A. Singla, and G. Soares. Generating high-precision feedback for programming syntax errors using large language models. *EDM 2023*, 2023.
 - [9] H. Joshi, J. C. Sanchez, S. Gulwani, V. Le, I. Radiček, and G. Verbruggen. Repair is nearly generation: Multilingual program repair with llms. In *Proceedings of the Thirty-Seventh AAAI Conference on Artificial Intelligence and Thirty-Fifth Conference on Innovative Applications of Artificial Intelligence and Thirteenth Symposium on Educational Advances in Artificial Intelligence*, AAAI'23/IAAI'23/EAAI'23. AAAI Press, 2023.
 - [10] H. Joshi, A. Ebenezer, J. Cambroner, S. Gulwani, A. Kanade, V. Le, I. Radiček, and G. Verbruggen. Flame: A small language model for spreadsheet formulas. *arXiv preprint arXiv:2301.13779 (AAAI 2024)*, 2023.
 - [11] A. D. Gordon, C. Negreanu, J. Cambroner, R. Chakravarthy, I. Drosos, H. Fang, B. Mitra, H. Richardson, A. Sarkar, S. Simmons, et al. Co-audit: tools to help humans double-check ai-generated content. *arXiv preprint arXiv:2310.01297 (PLATEAU 2024)*, 2023.
 - [12] J. Cambroner, S. Gulwani, V. Le, D. Perelman, A. Radhakrishna, C. Simon, and A. Tiwari. Flashfill++: Scaling programming by example by cutting to the chase. *Proceedings of the ACM on Programming Languages*, 7(POPL):952–981, 2023.
 - [13] J. Zhang, J. Cambroner, S. Gulwani, V. Le, R. Piskac, G. Soares, and G. Verbruggen. Repairing bugs in python assignments using large language models. *arXiv preprint arXiv:2209.14876 (to appear OOPSLA 2024)*, 2022.
 - [14] B. Wasti, J. P. Cambroner, B. Steiner, H. Leather, and A. Zlateski. Loopstack: a lightweight tensor algebra compiler stack. *arXiv preprint arXiv:2205.00618*, 2022.
 - [15] R. Bavishi, H. Joshi, J. Cambroner, A. Fariha, S. Gulwani, V. Le, I. Radiček, and A. Tiwari. Neurosymbolic repair for low-code formula languages. *Proc. ACM Program. Lang.*, 6(OOPSLA2), oct 2022.
 - [16] F. Zogaj, J. P. Cambroner, M. C. Rinard, and J. Cito. Doing more with less: characterizing dataset downsampling for automl. *Proceedings of the VLDB Endowment*, 14(11):2059–2072, 2021.
 - [17] M. Samak, J. P. Cambroner, and M. C. Rinard. Searching for replacement classes. *arXiv preprint arXiv:2110.05638*, 2021.
 - [18] T. H. Dang, J. P. Cambroner, and M. C. Rinard. Inferring drop-in binary parsers from program executions. *arXiv preprint arXiv:2104.09669*, 2021.
 - [19] L. Appelbaum, A. Berg, J. P. Cambroner, T. H. Y. Dang, C. C. Jin, L. Zhang, S. Kundrot, M. Palchuk, L. A. Evans, I. D. Kaplan, et al. Development of a pancreatic cancer prediction model using a multinational medical records database., 2021.

- [20] J. P. Cambronero, J. Cito, and M. C. Rinard. Ams: generating automl search spaces from weak specifications. In *Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, pages 763–774, 2020.
- [21] L. Appelbaum, J. P. Cambronero, J. P. Stevens, S. Horng, K. Pollick, G. Silva, S. Haneuse, G. Piatkowski, N. Benhaga, S. Duey, et al. Development and validation of a pancreatic cancer risk model for the general population using electronic health records: An observational study. *European Journal of Cancer*, 143:19–30, 2020.
- [22] L. Appelbaum, J. P. Cambronero, K. Pollick, G. Silva, J. P. Stevens, H. J. Mamon, I. D. Kaplan, and M. Rinard. Development and validation of a pancreatic cancer prediction model from electronic health records using machine learning., 2020.
- [23] J. P. Cambronero, J. Shen, J. Cito, E. Glassman, and M. Rinard. Characterizing developer use of automatically generated patches. In *2019 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*, pages 181–185. IEEE, 2019.
- [24] J. P. Cambronero and M. C. Rinard. Al: autogenerating supervised learning programs. *Proceedings of the ACM on Programming Languages*, 3(OOPSLA):1–28, 2019.
- [25] J. P. Cambronero, T. H. Dang, N. Vasilakis, J. Shen, J. Wu, and M. C. Rinard. Active learning for software engineering. In *Proceedings of the 2019 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software*, pages 62–78, 2019.
- [26] J. Cambronero, H. Li, S. Kim, K. Sen, and S. Chandra. When deep learning met code search. In *Proceedings of the 2019 27th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, pages 964–974, 2019.
- [27] J. Cambronero, P. Stanley-Marbell, and M. Rinard. Incremental color quantization for color-vision-deficient observers using mobile gaming data. *arXiv preprint arXiv:1803.08420*, 2018.
- [28] J. Cambronero, J. K. Feser, M. J. Smith, and S. Madden. Query optimization for dynamic imputation. *Proceedings of the VLDB Endowment*, 10(11):1310–1321, 2017.